

Assignment -4

| | |
|---------------------|-----------------------------|
| Team ID | PNT2022TMID15174 |
| Student Name | Charan C.D |
| Student Roll Number | 7179KCTKCTKCTKCTKCT19BEC130 |
| Maximum Marks | 2 Marks |

Input

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

Matplotlib is building the font cache; this may take a moment.

```
df=pd.read_csv("Mall_Customers.csv")
df
```

Output

| | CustomerID | Gender | Age | Annual Income (k\$) | Spending Score (1-100) |
|-----|------------|--------|-----|---------------------|------------------------|
| 0 | 1 | Male | 19 | 15 | 39 |
| 1 | 2 | Male | 21 | 15 | 81 |
| 2 | 3 | Female | 20 | 16 | 6 |
| 3 | 4 | Female | 23 | 16 | 77 |
| 4 | 5 | Female | 31 | 17 | 40 |
| ... | ... | ... | ... | ... | ... |
| 195 | 196 | Female | 35 | 120 | 79 |
| 196 | 197 | Female | 45 | 126 | 28 |
| 197 | 198 | Male | 32 | 126 | 74 |
| 198 | 199 | Male | 32 | 137 | 18 |
| 199 | 200 | Male | 30 | 137 | 83 |


200 rows × 5 columns

Input

```
df.drop(["CustomerID"],axis="columns",inplace=True)
```

```
sns.distplot(df.Age)
```


Output



Input

```
sns.displot(df["Annual Income (k$)"])
```


Output



Input

```
sns.lineplot(df.Age,df["Annual Income (k$)"])
```


Output



Input

```
sns.lineplot(df["Spending Score (1-100)"],df["Annual Income (k$)"])
```

Output



Input

```
df.describe()
```

Output

| | Age | Annual Income (k\$) | Spending Score (1-100) |
|-------|------------|---------------------|------------------------|
| count | 200.000000 | 200.000000 | 200.000000 |
| mean | 38.850000 | 60.560000 | 50.200000 |
| std | 13.969007 | 26.264721 | 25.823522 |
| min | 18.000000 | 15.000000 | 1.000000 |
| 25% | 28.750000 | 41.500000 | 34.750000 |
| 50% | 36.000000 | 61.500000 | 50.000000 |
| 75% | 49.000000 | 78.000000 | 73.000000 |
| max | 70.000000 | 137.000000 | 99.000000 |

Input

```
df.isnull().any()
```

Output

```
Gender          False
Age             False
Annual Income (k$)  False
Spending Score (1-100) False
dtype: bool
```

Input

```
df.isnull().sum()
```


Output

```
Gender          0
Age             0
Annual Income (k$)  0
Spending Score (1-100)  0
dtype: int64
```

Input

```
sns.boxplot(df.Age)
```


Output



Input

```
sns.boxplot(df["Annual Income (k$)"])
```

Output



Input

```
a=df["Annual Income (k$)"].quantile(0.99)  
a
```

Output


```
126.10999999999999
```

Input

```
df["Annual Income (k$)"]=np.where(df["Annual Income (k$)"]>=a,df["Annual Income  
(k$)"].median(),df["Annual Income (k$)"])
```

```
sns.boxplot(df["Annual Income (k$)"])
```


Output



Input

```
sns.boxplot(df["Spending Score (1-100)"])
```

Output



Input

```
from sklearn.preprocessing import LabelEncoder  
le=LabelEncoder()  
df.Gender=le.fit_transform(df.Gender)
```

```
df.head()
```

Output

| | Gender | Age | Annual Income (k\$) | Spending Score (1-100) |
|---|--------|-----|---------------------|------------------------|
| 0 | 1 | 19 | 15.0 | 39 |
| 1 | 1 | 21 | 15.0 | 81 |
| 2 | 0 | 20 | 16.0 | 6 |
| 3 | 0 | 23 | 16.0 | 77 |
| 4 | 0 | 31 | 17.0 | 40 |

Input

```
from sklearn.preprocessing import scale
df=pd.DataFrame(scale(df),columns=df.columns)
```

```
df.head()
```

Output

| | Gender | Age | Annual Income (k\$) | Spending Score (1-100) |
|---|-----------|-----------|---------------------|------------------------|
| 0 | 1.128152 | -1.424569 | -1.788777 | -0.434801 |
| 1 | 1.128152 | -1.281035 | -1.788777 | 1.195704 |
| 2 | -0.886405 | -1.352802 | -1.748853 | -1.715913 |
| 3 | -0.886405 | -1.137502 | -1.748853 | 1.040418 |
| 4 | -0.886405 | -0.563369 | -1.708930 | -0.395980 |

Input

```
from sklearn.cluster import KMeans
model=KMeans(n_clusters=2)
model
```

Output

```
KMeans(n_clusters=2)
```

Input

```
y_predicted=model.fit_predict(df)
y_predicted
```

Output

```
array([0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0,
       1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0,
       1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 1, 0,
       1, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0,
       0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1,
       1, 0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0,
       0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0,
       1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0,
       1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0,
       1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0,
       1, 0])
```

Input

```
df["clusters"]=y_predicted
df
```

Output

| | Gender | Age | Annual Income (k\$) | Spending Score (1-100) | clusters |
|-----|-----------|-----------|---------------------|------------------------|----------|
| 0 | 1.128152 | -1.424569 | -1.788777 | -0.434801 | 0 |
| 1 | 1.128152 | -1.281035 | -1.788777 | 1.195704 | 0 |
| 2 | -0.886405 | -1.352802 | -1.748853 | -1.715913 | 1 |
| 3 | -0.886405 | -1.137502 | -1.748853 | 1.040418 | 0 |
| 4 | -0.886405 | -0.563369 | -1.708930 | -0.395980 | 0 |
| ... | ... | ... | ... | ... | ... |
| 195 | -0.886405 | -0.276302 | 2.403201 | 1.118061 | 0 |
| 196 | -0.886405 | 0.441365 | 2.642742 | -0.861839 | 1 |
| 197 | 1.128152 | -0.491602 | 2.642742 | 0.923953 | 0 |
| 198 | 1.128152 | -0.491602 | 0.067670 | -1.250054 | 1 |
| 199 | 1.128152 | -0.635135 | 0.067670 | 1.273347 | 0 |

200 rows × 5 columns

Input

```
x=df.drop("clusters",axis="columns")
x
```

Output

| | Gender | Age | Annual Income (k\$) | Spending Score (1-100) |
|-----|-----------|-----------|---------------------|------------------------|
| 0 | 1.128152 | -1.424569 | -1.788777 | -0.434801 |
| 1 | 1.128152 | -1.281035 | -1.788777 | 1.195704 |
| 2 | -0.886405 | -1.352802 | -1.748853 | -1.715913 |
| 3 | -0.886405 | -1.137502 | -1.748853 | 1.040418 |
| 4 | -0.886405 | -0.563369 | -1.708930 | -0.395980 |
| ... | ... | ... | ... | ... |
| 195 | -0.886405 | -0.276302 | 2.403201 | 1.118061 |
| 196 | -0.886405 | 0.441365 | 2.642742 | -0.861839 |
| 197 | 1.128152 | -0.491602 | 2.642742 | 0.923953 |
| 198 | 1.128152 | -0.491602 | 0.067670 | -1.250054 |
| 199 | 1.128152 | -0.635135 | 0.067670 | 1.273347 |

200 rows × 4 columns

Input

```
y=df.clusters  
y
```

Output

```
0      0  
1      0  
2      1  
3      0  
4      0  
..  
195    0  
196    1  
197    0  
198    1  
199    0  
Name: clusters, Length: 200, dtype: int32
```

Input

```
from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=23)
```

```
x_train.shape
```

Output

```
(160, 4)
```

Input

```
x_test.shape
```

Output

```
(40, 4)
```

Input

```
lr.fit(x_train,y_train)
```

Output

```
LogisticRegression()
```

Input

```
lr.score(x_test,y_test)
```

Output

```
1.0
```

Input

```
lr.score(x_train,y_train)
```

Output

```
1.0
```


Input

```
from sklearn.metrics import confusion_matrix
y_pred = lr.predict(x_test)
cm = confusion_matrix(y_test, y_pred)
cm
```

Output


```
array([[19,  0],
       [ 0, 21]], dtype=int64)
```

Input

```
sns.heatmap(cm, annot=True)
plt.xlabel('Predicted')
plt.ylabel('Truth')
```

Output

```
Text(33.0, 0.5, 'Truth')
```



Input

```
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

Output

| | precision | recall | f1-score | support | |
|--------------|-----------|--------|----------|---------|----|
| | 0 | 1.00 | 1.00 | 1.00 | 19 |
| | 1 | 1.00 | 1.00 | 1.00 | 21 |
| accuracy | | | | 1.00 | 40 |
| macro avg | | 1.00 | 1.00 | 1.00 | 40 |
| weighted avg | | 1.00 | 1.00 | 1.00 | 40 |