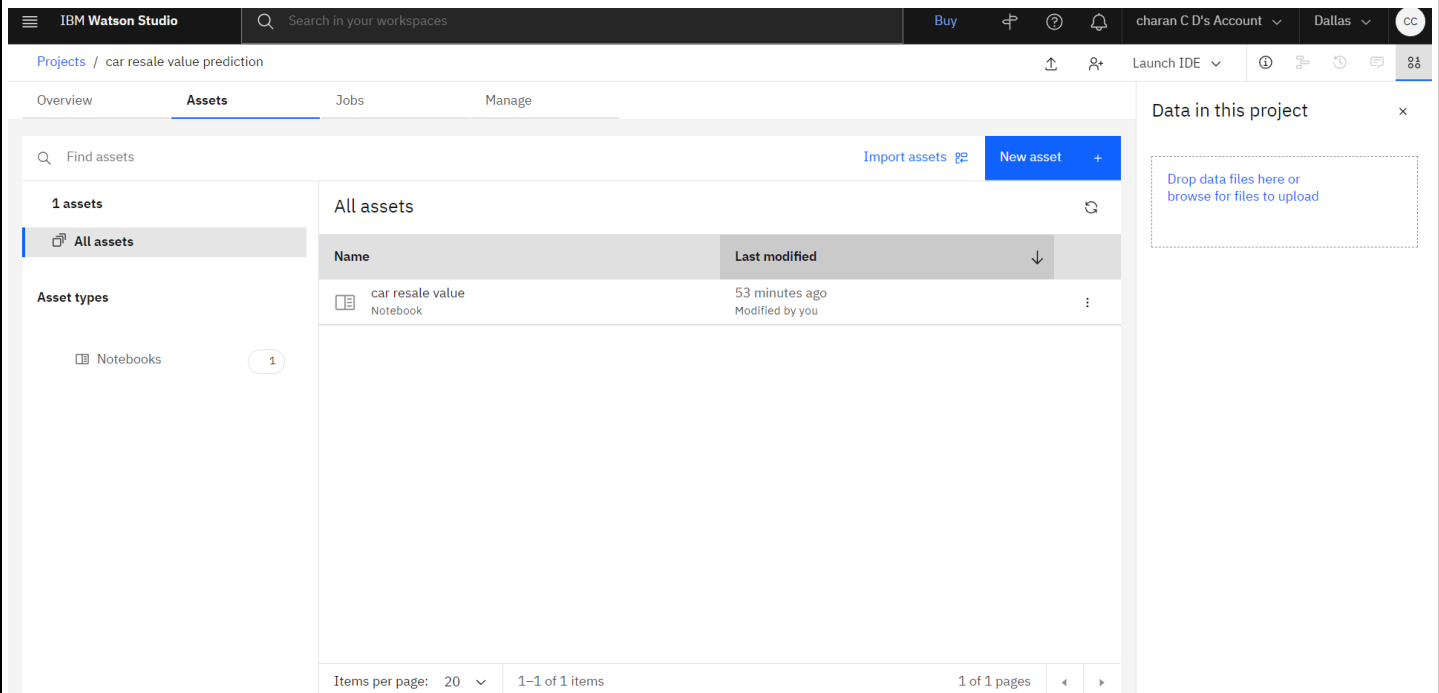


TRAIN THE ML MODEL ON IBM

Team ID	PNT2022TMID15174
Project Name	Car Resale value Prediction

TRAIN THE ML MODEL ON IBM



The screenshot shows the IBM Watson Studio interface. At the top, there's a navigation bar with 'IBM Watson Studio' and a search bar. Below it, the breadcrumb is 'Projects / car resale value prediction'. The main area has tabs for 'Overview', 'Assets', 'Jobs', and 'Manage'. The 'Assets' tab is active, showing a list of assets. On the left, there's a sidebar with '1 assets' and 'Asset types' (Notebooks). The main list shows one asset: 'car resale value Notebook', modified 53 minutes ago. On the right, there's a 'Data in this project' panel with a prompt to drop data files or browse for files to upload.

```
import pandas as pd
import numpy as np
import matplotlib as plt
from sklearn.preprocessing import LabelEncoder
import pickle
print("IMPORTED REQUIRED LIBRARIES")
# df = pd.read_csv("C:/Users/SUGARANJAN/Desktop/IBM/Data/autos.csv", header=0 , sep=',',
.encoding='Latin1',low_memory=False)
# df.head()
import os, types
import pandas as pd
from botocore.client import Config
import ibm_boto3
import io
def _iter(self): return 0

# @hidden_cell
# The following code accesses a file in your IBM Cloud Object Storage. It includes your credentials.
# You might want to remove those credentials before you share the notebook.
cos_client = ibm_boto3.client(service_name='s3',
    ibm_api_key_id='DT151-IL0017uhnUGwXyhG_Eort5gohoW6XJTNoT3RKK',
    ibm_auth_endpoint="https://iam.cloud.ibm.com/oidc/token",
```

```

config=Config(signature_version='oauth'),
endpoint_url='https://s3.private.us.cloud-object-storage.appdomain.cloud')

bucket = 'carresalevalueprediction-donotdelete-pr-yuhtzmzidi0ka1p'
object_key = 'autos.csv'

body = cos_client.get_object(Bucket=bucket,Key=object_key)
df = pd.read_csv((io.BytesIO(body['Body'].read())) , header=0 , sep=',',encoding='Latin1',low_memory=False)
df.head()
# df = pd.read_csv("C:/Users/SUGARANJAN/Desktop/IBM/Data/autos.csv", header=0 , sep=',',
,encoding='Latin1',low_memory=False)
# df.head()
import os, types
import pandas as pd
from botocore.client import Config
import ibm_boto3
import io
def _iter_(self): return 0

# @hidden_cell
# The following code accesses a file in your IBM Cloud Object Storage. It includes your credentials.
# You might want to remove those credentials before you share the notebook.
cos_client = ibm_boto3.client(service_name='s3',
    ibm_api_key_id='DT151-IL0017uhnUGwXyhG_Eort5gohoW6XJTNoT3RKK',
    ibm_auth_endpoint="https://iam.cloud.ibm.com/oidc/token",
    config=Config(signature_version='oauth'),
    endpoint_url='https://s3.private.us.cloud-object-storage.appdomain.cloud')

bucket = 'carresalevalueprediction-donotdelete-pr-yuhtzmzidi0ka1p'
object_key = 'autos.csv'

body = cos_client.get_object(Bucket=bucket,Key=object_key)
df = pd.read_csv((io.BytesIO(body['Body'].read())) , header=0 , sep=',',encoding='Latin1',low_memory=False)
df.head()
print(df.seller.value_counts())
df[df.seller != 'gewerblich']
df=df.drop('seller',axis=1)

print(df.offerType.value_counts())
df[df.offerType != 'Gesuch']
df=df.drop('offerType',axis=1)
print(df.shape)
df=df[(df.powerPS>50) & (df.powerPS<900)]
print(df.shape)
df=df[(df.yearOfRegistration>=1950)&(df.yearOfRegistration<2022)]
print(df.shape)
df.drop(['name','abtest','dateCrawled','nrOfPictures','lastSeen','postalCode','dateCreated'], axis='columns',inplace=True)
new_df=df.copy()
new_df=new_df.drop_duplicates(['price','vehicleType','yearOfRegistration','gearbox','powerPS','model','kilometer','monthOfRegistration','fuelType','notRepairedDamage'])
new_df.gearbox.replace(('manuell','automatik'),('manual','automatic'),inplace=True)
new_df.fuelType.replace(('benzin','andere','elektro'),('petrol','others','electric'),inplace=True)
new_df.vehicleType.replace(('kleinwagen','cabrio','kombi','andere'),('small car','convertible','combination','others'),inplace=True)

```

```

new_df.notRepairedDamage.replace(('ja','nein'),('Yes','No'),inplace=True)
new_df=new_df[(new_df.price>=100)&(new_df.price<=150000)]

new_df['notRepairedDamage'].fillna(value='not-declared',inplace=True)
new_df['fuelType'].fillna(value='not-declared',inplace=True)
new_df['gearbox'].fillna(value='not-declared',inplace=True)
new_df['vehicleType'].fillna(value='not-declared',inplace=True)
new_df['model'].fillna(value='not-declared',inplace=True)
from ibm_watson_machine_learning import APIClient
wml_credentials={
    "url":"https://us-south.ml.cloud.ibm.com",
    "apikey":"hEAn_mcoP3u_-ZjagjeqlxDayqUiETpYVYWdR1OLKAby"
}
client =APIClient(wml_credentials)
def guide_from_space_name(client, space_name):
    space = client.spaces.get_details()
    # print(space)
    return(next(item for item in space['resources'] if item['entity']['name']==space_name)['metadata']['id'])
space_uid=guide_from_space_name(client,'CAR')
print("Space UID"+ space_uid)
client.set.default_space(space_uid)
client.software_specifications.list()
software_spec_uid = client.software_specifications.get_uid_by_name("runtime-22.1-py3.9")
software_spec_uid
print(new_df)
labels=['gearbox','notRepairedDamage','model','brand','fuelType','vehicleType']

mapper={}
for i in labels:
    mapper[i]=LabelEncoder()
    mapper[i].fit(new_df[i])
    tr=mapper[i].transform(new_df[i])
    np.save(str('classes'+i+'.npy'),mapper[i].classes_)
    print(i,":",mapper[i])
    new_df.loc[:, i+ '_labels']=pd.Series(tr,index=new_df.index)

labeled = new_df[['price','yearOfRegistration','powerPS','kilometer','monthOfRegistration']+['x+'+_labels" for x in
labels]]
print(labeled.columns)
Y=labeled.iloc[:,0].values
X=labeled.iloc[:,1:].values

Y=Y.reshape(-1,1)
from sklearn.model_selection import cross_val_score,train_test_split
X_train , X_test, Y_train , Y_test = train_test_split(X,Y,test_size=0.3,random_state=3)
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score
regressor = RandomForestRegressor(n_estimators = 1000,max_depth = 10,random_state = 34)

regressor.fit(X_train, np.ravel(Y_train,order='C'))
y_pred = regressor.predict(X_test)
print(r2_score(Y_test,y_pred))
filename='resale_model.sav'
pickle.dump(regressor,open(filename,'wb'))

```

```

model_details = client.repository.store_model(model=regressor,meta_props={
    client.repository.ModelMetaNames.NAME: "resale_model",
    client.repository.ModelMetaNames.SOFTWARE_SPEC_UID: software_spec_uid,
    client.repository.ModelMetaNames.TYPE: "scikit-learn_1.0"
})
model_id = client.repository.get_model_id(model_details)
model_id
X_train[0]
regressor.predict([[2012.0, 179.0, '1500000', 12.0, 0, 0, 30, 1, 1, 4]])

```

IBM Watson Studio interface showing a Jupyter notebook. The top bar includes the IBM Watson Studio logo, a search bar, and user account information (charan C D's Account, Dallas). The notebook is titled "car resale value prediction" and "car resale value". The code in the notebook is as follows:

```

In [1]: import pandas as pd
import numpy as np
import matplotlib as plt
from sklearn.preprocessing import LabelEncoder
import pickle
print("IMPORTED REQUIRED LIBRARIES")

IMPORTED REQUIRED LIBRARIES

In [9]: # df = pd.read_csv("C:\Users\Happy\Downloads\IBM Project\datasheets\Data\autos.csv", header=0, sep=',', encoding='Latin1', low_memory=False)
# df.head()
import os, types
import pandas as pd
from botocore.client import Config
import ibm_boto3
import io
def iter (self): return 0

# @hidden_cell
# The following code accesses a file in your IBM Cloud Object Storage. It includes your credentials.
# You might want to remove those credentials before you share the notebook.
cos_client = ibm_boto3.client(service_name='s3',
                              ibm_api_key_id='DT151-1L0017uhnuGwXyHG_Eort5gohom6XJTMoT3Rkk',
                              ibm_auth_endpoint="https://iam.cloud.ibm.com/oidc/token",
                              config=Config(signature_version='oauth'),
                              endpoint_url='https://s3.private.us.cloud-object-storage.appdomain.cloud')

bucket = 'carresalevalueprediction-donotdelete-pr-yuhtmzidi0ka1p'
object_key = 'autos.csv'

body = cos_client.get_object(Bucket=bucket,Key=object_key)
df = pd.read_csv(io.BytesIO(body['Body'].read()), header=0, sep=',', encoding='Latin1', low_memory=False)
df.head()

```

IBM Watson Studio interface showing a Jupyter notebook. The top bar includes the IBM Watson Studio logo, a search bar, and user account information (charan C D's Account, Dallas). The notebook is titled "car resale value prediction" and "car resale value". The code in the notebook is as follows:

```

In [5]: print(df.seller.value_counts())
df[df.seller != 'gewerblich']
df=df.drop('seller',axis=1)

print(df.offerType.value_counts())
df[df.offerType != 'gesucht']
df=df.drop('offerType',axis=1)

privat      371534
gewerblich    3
golf         1
Name: seller, dtype: int64
Angebot      371535
Gesuch       12
150000       1
Name: offerType, dtype: int64

In [6]: print(df.shape)
df=df[(df.powerPS>50) & (df.powerPS<900)]
print(df.shape)
df=df[(df.yearOfRegistration>=1950)&(df.yearOfRegistration<2022)]
print(df.shape)

(371539, 18)
(319717, 18)
(319649, 18)

In [7]: df.drop(['name','abtest','dateCrawled','nrOfPictures','lastSeen','postalCode','dateCreated'], axis='columns',inplace=True)

In [8]: new_df=df.copy()
new_df=new_df.drop_duplicates(['price','vehicleType','yearOfRegistration','gearbox','powerPS','model','kilometer','monthOfRegistration','fuelType','notRepairedDamage'])

In [9]: new_df.gearbox.replace({'manuell','automatik'},('manual','automatic'),inplace=True)
new_df.fuelType.replace({'benzin','andere','elektro'},('petrol','others','electric'),inplace=True)
new_df.vehicleType.replace({'kleinwagen','cabrio','kombi','andere'},('small car','convertible','combination','others'),inplace=True)
new_df.notRepairedDamage.replace({'ja','nein'},('Yes','No'),inplace=True)

```

```
IBM Watson Studio Search in your workspaces Buy charan C D's Account Dallas CC

Projects / car resale value prediction / car resale value

File Edit View Insert Cell Kernel Help Trusted | Python 3.9

In [7]: df.drop(['name', 'abtest', 'dateCrawled', 'numberOfPictures', 'lastSeen', 'postalCode', 'dateCreated'], axis='columns', inplace=True)

In [8]: new_df=df.copy()
new_df=new_df.drop_duplicates(['price', 'vehicleType', 'yearOfRegistration', 'gearbox', 'powerPS', 'model', 'kilometer', 'monthOfRegistration', 'fuelType', 'notRepairedDamage'])

In [9]: new_df.gearbox.replace({'manuel', 'automatik'}, {'manual', 'automatic'}, inplace=True)
new_df.fuelType.replace({'benzin', 'andere', 'elektro'}, {'petrol', 'others', 'electric'}, inplace=True)
new_df.vehicleType.replace({'kleinwagen', 'cabrio', 'kombi', 'andere'}, {'small car', 'convertible', 'combination', 'others'}, inplace=True)
new_df.notRepairedDamage.replace({'ja', 'nein'}, {'Yes', 'No'}, inplace=True)

In [10]: new_df=new_df[(new_df.price>=100)&(new_df.price<=150000)]

new_df['notRepairedDamage'].fillna(value='not-declared', inplace=True)
new_df['fuelType'].fillna(value='not-declared', inplace=True)
new_df['gearbox'].fillna(value='not-declared', inplace=True)
new_df['vehicleType'].fillna(value='not-declared', inplace=True)
new_df['model'].fillna(value='not-declared', inplace=True)

In [11]: from ibm_watson_machine_learning import APIClient
uml_credentials={
    "url": "https://us-south.ml.cloud.ibm.com",
    "apikey": "hEAn_mcoP3u_-ZjagjeqlxDayQUIETpYVYdR10LEAby"
}
client =APIClient(uml_credentials)

In [12]: def guide_from_space_name(client, space_name):
    space = client.spaces.get_details()
    # print(space)
    return(next(item for item in space['resources'] if item['entity']['name']==space_name)['metadata']['id'])

In [13]: space_uid=guide_from_space_name(client, 'CAR')
print("Space UID="+ space_uid)

Space UIDbe467bbb-03a2-40e7-bf5a-91836e346951

In [14]: client.set.default_space(space_uid)

Out[14]: 'SUCCESS'
```

```
IBM Watson Studio Search in your workspaces Buy charan C D's Account Dallas CC

Projects / car resale value prediction / car resale value

File Edit View Insert Cell Kernel Help Trusted | Python 3.9

In [15]: client.software_specifications.list()

-----
NAME ASSET_ID TYPE
default_py3.6 0062b8c9-8b7d-44a0-a0b9-46c416adcbd9 base
kernel-spark3.2-scala2.12 020d69ce-7ac1-5e68-ac1a-31189867356a base
pytorch-onnx_1.3-py3.7-edt 069ea134-3346-5748-b513-49120e15d288 base
scikit-learn_0.20-py3.6 09c5a1d0-9c1e-4473-a344-cb7b665ff687 base
spark-mllib_3.0-scala_2.12 09f4cffe-90a7-5099-b9ed-1ef348aebdee base
pytorch-onnx_rt22.1-py3.9 0b848dd4-e581-5599-be41-b5f6fcccc6471 base
ai-function_0.1-py3.6 0cdb0f1e-5376-4f4d-92dd-da3b69a9bda base
shiny-r3.6 0e6e79df-875e-4f24-8ae9-62dc2148306 base
tensorflow_2.4-py3.7-horovod 1092590a-307d-563d-9b62-4eb7d64b3f22 base
pytorch_1.1-py3.6 10ac12d6-6b30-4ccd-8392-3e922c096a92 base
tensorflow_1.15-py3.6-ddl 111e41b3-de2d-5422-a4d6-bf776828c4b7 base
autoai-kb_rt22.2-py3.10 125b6d9a-5b1f-5e8d-972a-b251688ccf40 base
runtime-22.1-py3.9 12b83a17-24d8-5082-900f-0ab31fbfd3cb base
scikit-learn_0.22-py3.6 154010fa-5b3b-4ac1-82af-4d5ee5abbc85 base
default_r3.6 1b70aee1-ab34-4b87-8aa0-a4a3c8296a36 base
pytorch-onnx_1.3-py3.6 1bc6029a-cc97-56da-b8e0-39c3880dbbe7 base
kernel-spark3.3-r3.6 1c9e5454-f216-59dd-a20e-474a5cdf5988 base
pytorch-onnx_rt22.1-py3.9-edt 1d362186-7ad5-5b59-8b6c-9d0880bde37f base
tensorflow_2.1-py3.6 1eb25b84-d6ed-5dde-b6a5-3fbfd1665666 base
spark-mllib_3.2 20047f72-0a98-58c7-9ff5-a77b012eb8f5 base
tensorflow_2.4-py3.8-horovod 217c16f6-178f-56bf-824a-b19f20564c49 base
runtime-22.1-py3.9-cuda 26215f05-08c3-5a41-a1b0-da66306ce658 base
do_py3.8 295addb5-9ef9-547e-9bfa-92ae3563e720 base
autoai-ts_3.0-py3.8 2aa0c932-798f-5ae9-abd6-15e0c2402fb5 base
tensorflow_1.15-py3.6 2b73a275-7cbf-420b-a912-eae7f436e0bc base
kernel-spark3.3-py3.9 2b7961e2-e3b1-5a8c-a491-482c8368839a base
pytorch_1.2-py3.6 2c8ef57d-2687-4b7d-acce-01f94976dac1 base
spark-mllib_2.3 2e517f00-bca0-4b0d-88dc-5c6791338875 base
pytorch-onnx_1.1-py3.6-edt 32983cea-3f32-4400-8965-dde874a8d67e base
spark-mllib_3.0-py37 36507ebe-8770-55ba-ab2a-eafe787600e9 base
spark-mllib_2.4 390d21f8-e58b-4fac-9c55-d7ceda621326 base
autoai-ts_rt22.2-py3.10 396b2e83-0953-5b86-9a55-7ce1628a406f base
xgboost_0.82-py3.6 39e31acd-5f30-41dc-ae44-80233c80306e base
pytorch-onnx_1.2-py3.6-edt 40589d0e-7019-4e28-8daa-fb03b6f4fe12 base
pytorch-onnx_rt22.2-py3.10 40e73f55-783a-5535-b3fa-0c8b94291431 base
```

```
IBM Watson Studio Search in your workspaces Buy charan C D's Account Dallas CC

Projects / car resale value prediction / car resale value

File Edit View Insert Cell Kernel Help Trusted | Python 3.9

In [18]: labels=['gearbox','notRepairedDamage','model','brand','fuelType','vehicleType']

mapper={}
for i in labels:
    mapper[i]=LabelEncoder()
    mapper[i].fit(new_df[i])
    tr=mapper[i].transform(new_df[i])
    np.save(str('classes'+i+'.npy'),mapper[i].classes_)
    print(i,":",mapper[i])
    new_df.loc[:, i+ '_labels']=pd.Series(tr,index=new_df.index)

labeled = new_df[['price','yearOfRegistration','powerPS','kilometer','monthOfRegistration']+['x'+ "_labels" for x in labels]]
print(labeled.columns)

gearbox : LabelEncoder()
notRepairedDamage : LabelEncoder()
model : LabelEncoder()
brand : LabelEncoder()
fuelType : LabelEncoder()
vehicleType : LabelEncoder()
Index(['price', 'yearOfRegistration', 'powerPS', 'kilometer',
       'monthOfRegistration', 'gearbox_labels', 'notRepairedDamage_labels',
       'model_labels', 'brand_labels', 'fuelType_labels',
       'vehicleType_labels'],
      dtype='object')

In [19]: Y=labeled.iloc[:,0].values
X=labeled.iloc[:,1:].values

Y=Y.reshape(-1,1)

In [20]: from sklearn.model_selection import cross_val_score,train_test_split
X_train , X_test, Y_train , Y_test = train_test_split(X,Y,test_size=0.3,random_state=3)

In [21]: from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score
```

```
IBM Watson Studio Search in your workspaces Buy charan C D's Account Dallas CC

Projects / car resale value prediction / car resale value

File Edit View Insert Cell Kernel Help Trusted | Python 3.9

In [16]: software_spec_uid = client.software_specifications.get_uid_by_name("runtime-22.1-py3.9")
software_spec_uid

Out[16]: '12b83a17-24d8-5082-900f-0ab31fbfd3cb'

In [17]: print(new_df)

   price  vehicleType  yearOfRegistration  gearbox  powerPS  \
1  18300.0         coupe             2011.0   manual    190.0
2   9800.0          suv              2004.0  automatic   163.0
3   1500.0      samll car             2001.0   manual    75.0
4   3600.0      samll car             2008.0   manual    69.0
5    650.0     limousine             1995.0   manual   102.0
...     ...         ...                 ...     ...     ...
371531  3200.0     limousine             2004.0   manual   215.0
371535  1199.0  convertible             2000.0  automatic   101.0
371536  9200.0         bus              1996.0   manual   102.0
371537  3400.0  combination             2002.0   manual   100.0
371538 28990.0     limousine             2013.0   manual   320.0

   model  kilometer  monthOfRegistration  fuelType  brand  \
1  not-declared   125000                5.0   diesel   audi
2    grand       125000                8.0   diesel   jeep
3    golf       150000                6.0   petrol  volkswagen
4    fabia       90000                7.0   diesel   skoda
5     3er       150000               10.0   petrol    bmw
...     ...         ...                 ...     ...     ...
371531    leon       150000                5.0   petrol   seat
371535  fortwo       125000                3.0   petrol  smart
371536  transporter  150000                3.0   diesel  volkswagen
371537    golf       150000                6.0   diesel  volkswagen
371538   m_reihe       50000                8.0   petrol    bmw

   notRepairedDamage
1                Yes
2          not-declared
3                  Na
```

IBM Watson Studio

Search in your workspaces

Buy

charan C D's Account

Dallas

CC

Projects / car resale value prediction / car resale value

File Edit View Insert Cell Kernel Help

Trusted | Python 3.9

```
regressor = RandomForestRegressor(n_estimators = 1000,max_depth = 10,random_state = 34)
regressor.fit(X_train, np.ravel(Y_train,order='C'))

Out[21]: RandomForestRegressor(max_depth=10, n_estimators=1000, random_state=34)

In [22]: y_pred = regressor.predict(X_test)
print(r2_score(Y_test,y_pred))

0.8310350387286918

In [23]: filename='resale_model.sav'
pickle.dump(regressor,open(filename,'wb'))

In [24]: model_details = client.repository.store_model(model=regressor,meta_props={
    client.repository.ModelMetaNames.NAME: "resale_model",
    client.repository.ModelMetaNames.SOFTWARE_SPEC_UID: software_spec_uid,
    client.repository.ModelMetaNames.TYPE: "scikit-learn_1.0"
})
model_id = client.repository.get_model_id(model_details)

In [25]: model_id

Out[25]: 'cd8479e0-66e4-454e-aecce-4824fe9d71bd'

In [26]: X_train[0:]

Out[26]: array([[2005.0, 179.0, '150000', ..., 1, 1, 4],
 [1997.0, 60.0, '150000', ..., 38, 7, 4],
 [2003.0, 170.0, '150000', ..., 2, 7, 1],
 ...,
 [2009.0, 174.0, '125000', ..., 25, 7, 7],
 [2000.0, 136.0, '150000', ..., 20, 7, 1],
 [2013.0, 170.0, '40000', ..., 1, 7, 8]], dtype=object)

In [27]: regressor.predict([[2012.0, 179.0, '1500000', 12.0, 0, 0, 30, 1, 1, 4]])

Out[27]: array([18618.270101351])
```

